



Προγραμματισμός Συστημάτων Υψηλών Επιδόσεων (HY421 / MΔΕ646)

Εισαγωγή στο Μοντέλο Προγραμματισμού CUDA

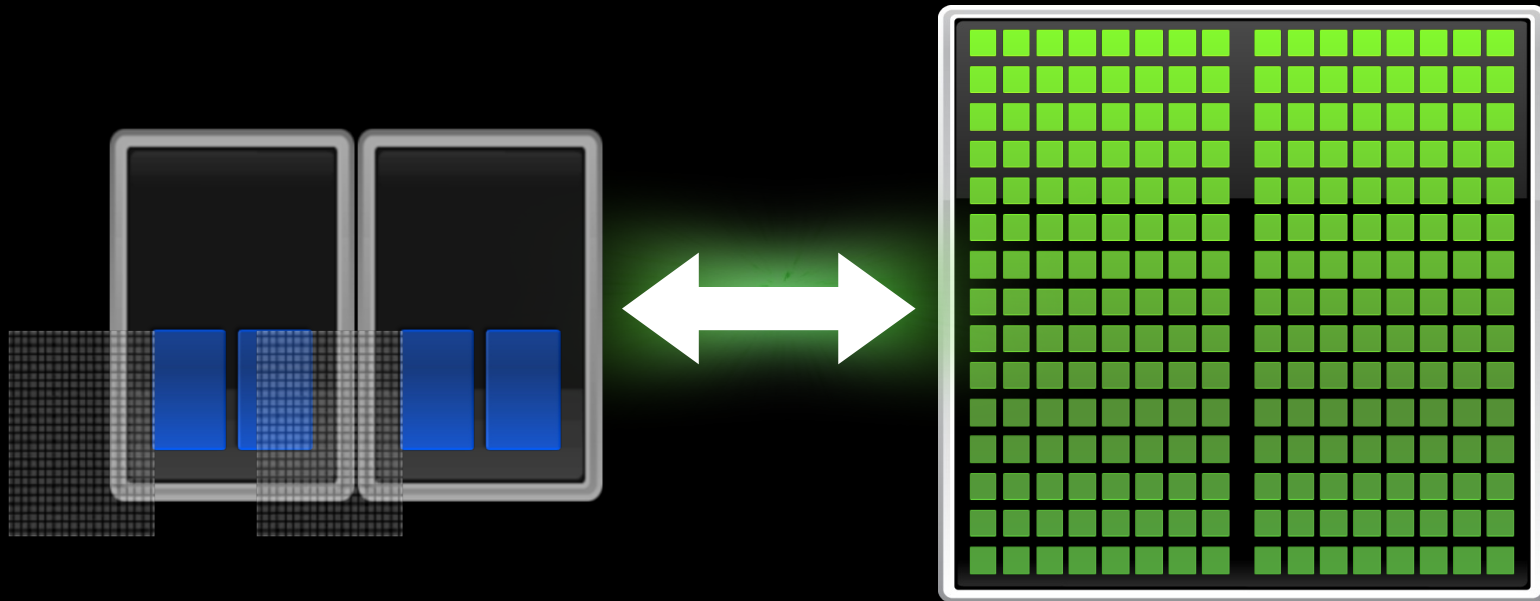
Heterogeneous Computing



2

Multicore CPU

Manycore GPU



C for CUDA



3

- Philosophy: provide minimal set of extensions necessary to expose power

- Function qualifiers:

```
__global__ void my_kernel() { }  
__device__ float my_device_func() { }
```

- Variable qualifiers:

```
__constant__ float my_constant_array[32];  
__shared__ float my_shared_array[32];
```

- Execution configuration:

```
dim3 grid_dim(100, 50); // 5000 thread blocks  
dim3 block_dim(4, 8, 8); // 256 threads per block  
my_kernel <<< grid_dim, block_dim >>> (...); // Launch kernel
```

- Built-in variables and functions valid in device code:

```
dim3 blockDim; // Grid dimension  
dim3 blockDim; // Block dimension  
dim3 blockIdx; // Block index  
dim3 threadIdx; // Thread index  
void __syncthreads(); // Thread synchronization
```

Example: vector_addition



4

Device Code

```
// compute vector sum c = a + b
// each thread performs one pair-wise addition
__global__ void vector_add(float* A, float* B, float* C)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    C[i] = A[i] + B[i];
}
```

```
int main()
{
    // elided initialization code
    ...
    // Run N/256 blocks of 256 threads each
    vector_add<<< N/256, 256>>>>(d_A, d_B, d_C);
}
```

Example: vector_addition



5

```
// compute vector sum c = a + b
// each thread performs one pair-wise addition
__global__ void vector_add(float* A, float* B, float* C)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    C[i] = A[i] + B[i];
}
```

Host Code

```
int main()
{
    // elided initialization code
    ...
    // launch N/256 blocks of 256 threads each
    vector_add<<< N/256, 256>>>>(d_A, d_B, d_C);
}
```

Example: Initialization code for vector_addition



6

```
// allocate and initialize host (CPU) memory  
float *h_A = ..., *h_B = ...;
```

```
// allocate device (GPU) memory  
float *d_A, *d_B, *d_C;
```

```
cudaMalloc( (void**) &d_A, N * sizeof(float));  
cudaMalloc( (void**) &d_B, N * sizeof(float));  
cudaMalloc( (void**) &d_C, N * sizeof(float));
```

```
// copy host memory to device
```

```
cudaMemcpy( d_A, h_A, N * sizeof(float),  
            cudaMemcpyHostToDevice );  
cudaMemcpy( d_B, h_B, N * sizeof(float),  
            cudaMemcpyHostToDevice );
```

```
// launch N/256 blocks of 256 threads each  
vector_add<<<N/256, 256>>>(d_A, d_B, d_C);
```



BASIC KERNELS AND EXECUTION ON GPU

CUDA Programming Model



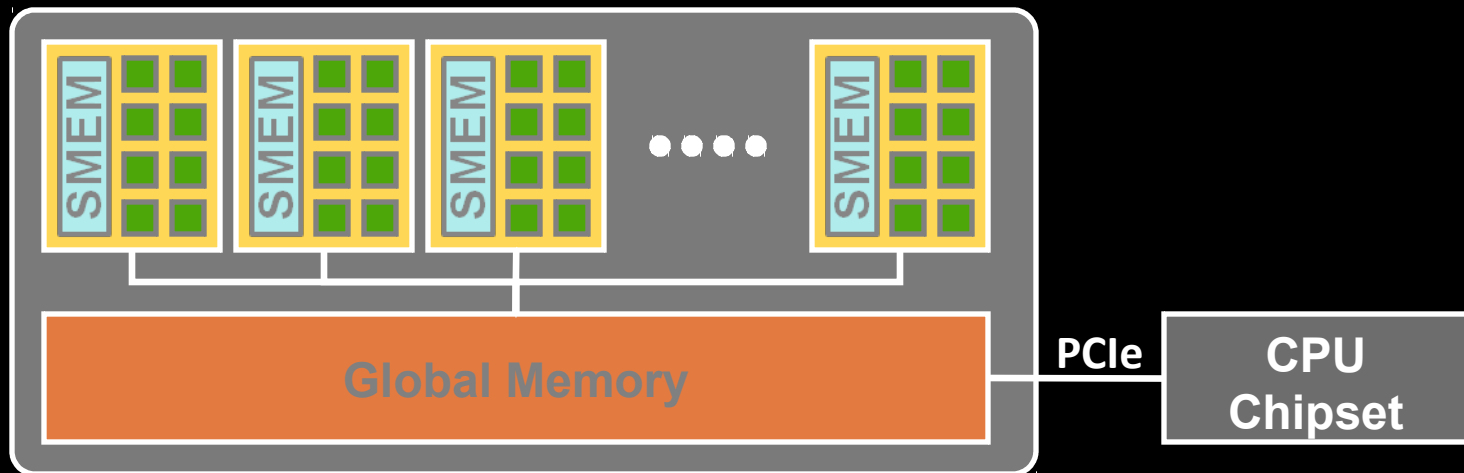
8

- **Parallel code (kernel) is launched and executed on a device by many threads**
- **Launches are hierarchical**
- **Threads are grouped into blocks**
- **Blocks are grouped into grids**
- **Familiar serial code is written for a thread**
- **Each thread is free to execute a unique code path**
- **Built-in thread and block ID variables**

High Level View



9



Blocks of threads run on an SM



10

Streaming Processor



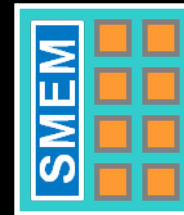
Thread

Registers

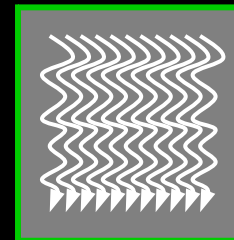


Memory

Streaming Multiprocessor



Threadblock



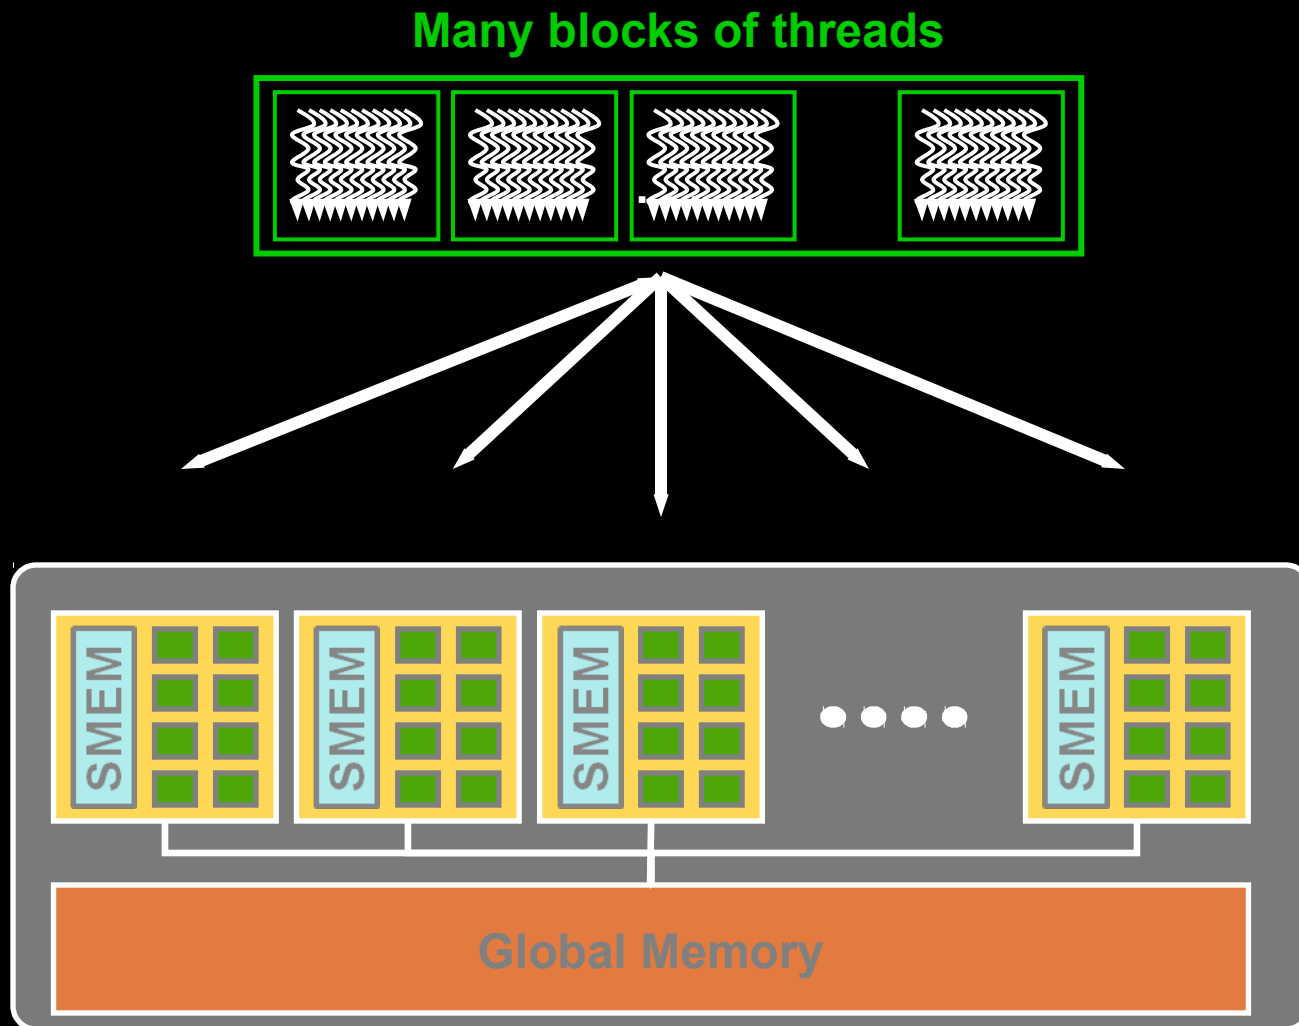
Per-block
Shared
Memory

Memory

Whole grid runs on GPU



11



Thread Hierarchy



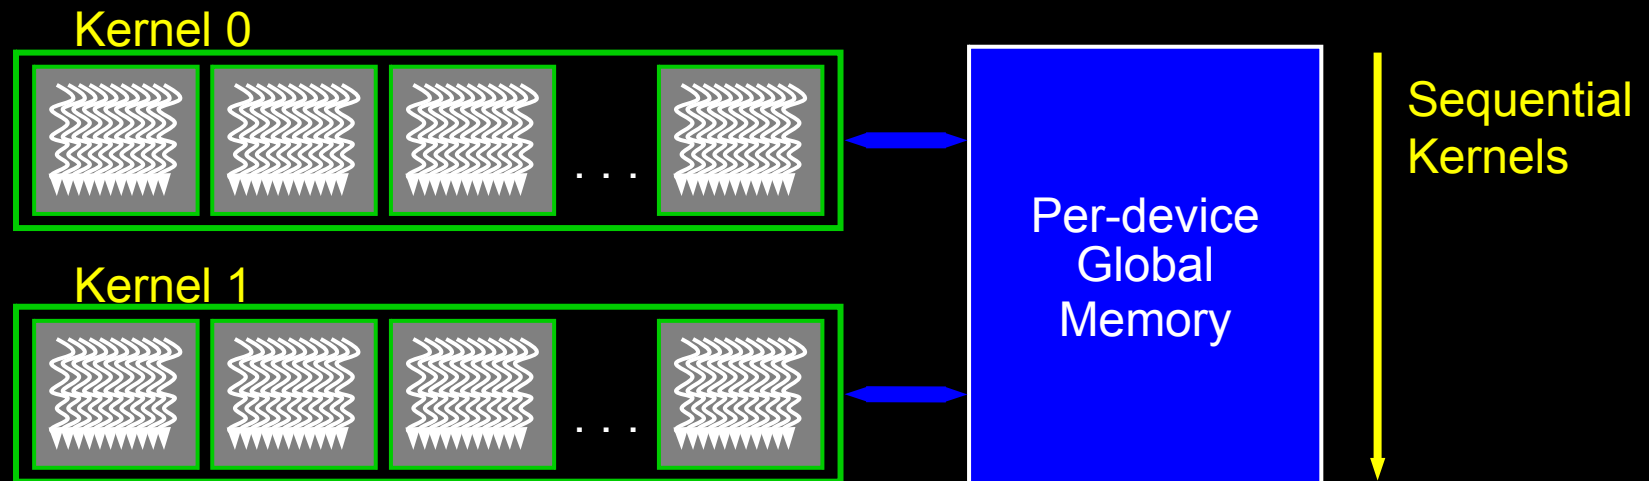
12

- **Threads launched for a parallel section are partitioned into thread blocks**
- **Grid = all blocks for a given launch**
- **Thread block is a group of threads that can:**
 - **Synchronize their execution**
 - **Communicate via shared memory**

Memory Model



13



Memory Model



14



Example: Vector Addition Kernel



15

Device Code

```
// Compute vector sum C = A+B
// Each thread performs one pair-wise addition
__global__ void vecAdd(float* A, float* B, float*
C)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    C[i] = A[i] + B[i];
}
```

```
int main()
{
    // Run grid of N/256 blocks of 256 threads
    each
```

```
vecAdd<<< N/256, 256>>>>(d_A, d_B, d_C);
```

Example: Vector Addition Kernel



16

```
// Compute vector sum C = A+B
// Each thread performs one pair-wise addition
__global__ void vecAdd(float* A, float* B, float*
C)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    C[i] = A[i] + B[i];
}
```

Host Code

```
int main()
{
    // Run grid of N/256 blocks of 256 threads
    each
```

```
vecAdd<<< N/256, 256>>>>(d_A, d_B, d_C);
```


Example: Host code for vecAdd



17

```
// allocate and initialize host (CPU) memory  
float *h_A = ..., *h_B = ...; *h_C = ... (empty)
```

```
// allocate device (GPU) memory
```

```
float *d_A, *d_B, *d_C;  
cudaMalloc( (void**) &d_A, N * sizeof(float));  
cudaMalloc( (void**) &d_B, N * sizeof(float));  
cudaMalloc( (void**) &d_C, N * sizeof(float));
```

```
// copy host memory to device
```

```
cudaMemcpy( d_A, h_A, N * sizeof(float),  
            cudaMemcpyHostToDevice );  
cudaMemcpy( d_B, h_B, N * sizeof(float),  
            cudaMemcpyHostToDevice );
```

```
// execute grid of N/256 blocks of 256 threads each  
vecAdd<<<N/256, 256>>>(d_A, d_B, d_C);
```

Example: Host code for vecAdd (2)



18

```
// execute grid of N/256 blocks of 256 threads each  
vecAdd<<<N/256, 256>>>(d_A, d_B, d_C);
```

```
// copy result back to host memory
```

```
cudaMemcpy( h_C, d_C, N * sizeof(float),  
            cudaMemcpyDeviceToHost );
```

```
// do something with the result...
```

```
// free device (CPU) memory
```

```
cudaFree(d_A);
```

```
cudaFree(d_B);
```

```
cudaFree(d_C);
```

Kernel Variations and Output



19

```
__global__ void kernel( int *a )  
{  
    int idx = blockIdx.x*blockDim.x + threadIdx.x;  
    a[idx] = 7;  
}
```

Output: 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7

```
__global__ void kernel( int *a )  
{  
    int idx = blockIdx.x*blockDim.x + threadIdx.x;  
    a[idx] = blockIdx.x;  
}
```

Output: 0 0 0 0 1 1 1 1 2 2 2 2 3 3 3 3

```
__global__ void kernel( int *a )  
{  
    int idx = blockIdx.x*blockDim.x + threadIdx.x;  
    a[idx] = threadIdx.x;  
}
```

Output: 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3

Code executed on GPU



20

- **C/C++ with some restrictions:**
- Can only access GPU memory (well...)
- No variable number of arguments
- No static variables
- No dynamic polymorphism
- **Must be declared with a qualifier:**
- **__global__** : launched by CPU,
cannot be called from GPU must return void
- **__device__** : called from other GPU functions,
cannot be called by the CPU
- **__host__** : can be called by CPU
- **__host__** and **__device__** qualifiers can be combined
- sample use: overloading operators

Memory Spaces



21

- **CPU and GPU have separate memory spaces**
- Data is moved across PCIe bus
- Use functions to allocate/set/copy memory on GPU
- Very similar to corresponding C functions
- **Pointers are just addresses**
- Can't tell from the pointer value whether the address is on CPU or GPU
- Must exercise care when dereferencing:
- Dereferencing CPU pointer on GPU will likely crash
 - Except on really new architectures and versions of CUDA (unified address space)
- Same for vice versa

GPU Memory Allocation / Release



22

- Host (CPU) manages device (GPU) memory:
- `cudaMalloc (void ** pointer, size_t nbytes)`
- `cudaMemset (void * pointer, int value, size_t count)`
- `cudaFree (void* pointer)`

```
int n = 1024;
```

```
int nbytes = 1024*sizeof(int);
```

```
int * d_a = 0;
```

```
cudaMalloc( (void**)&d_a, nbytes );
```

```
cudaMemset( d_a, 0, nbytes);
```

```
cudaFree(d_a);
```

Data Copies



23

- **cudaMemcpy(void *dst, void *src, size_t nbytes, enum cudaMemcpyKind direction);**
- returns after the copy is complete
- blocks CPU thread until all bytes have been copied
- doesn't start copying until previous CUDA calls complete
- **enum cudaMemcpyKind**
- **cudaMemcpyHostToDevice**
- **cudaMemcpyDeviceToHost**
- **cudaMemcpyDeviceToDevice**
- **Non-blocking copies are also available**

Dummy Code Walkthrough 1



24

```
// walkthrough1.cu  
#include <stdio.h>
```

```
int main()  
{
```

```
    int dimx = 16;  
    int num_bytes = dimx*sizeof(int);
```

```
    int *d_a=0, *h_a=0; // device and host pointers
```


Dummy Code Walkthrough 1



25

```
// walkthrough1.cu
#include <stdio.h>

int main()
{
    int dimx = 16;
    int num_bytes = dimx*sizeof(int);

    int *d_a=0, *h_a=0; // device and host pointers

    h_a = (int*)malloc(num_bytes);
    cudaMalloc( (void**)&d_a, num_bytes );

    if( h_a == NULL || d_a == NULL )
    {
        printf("couldn't allocate memory\n");
        return 1;
    }
}
```

Dummy Code Walkthrough 1



26

```
// walkthrough1.cu  
#include <stdio.h>
```

```
int main()  
{  
    int dimx = 16;  
    int num_bytes = dimx*sizeof(int);  
  
    int *d_a=NULL, *h_a=NULL // device and host pointers  
  
    h_a = (int*)malloc(num_bytes);  
    cudaMalloc( (void**)&d_a, num_bytes );  
  
    if(h_a == NULL || d_a == NULL )  
    {  
        printf("couldn't allocate memory\n");  
        return 1;  
    }  
  
    cudaMemset( d_a, 0, num_bytes );  
    // do something on the device to fill the buffer  
    cudaMemcpy( h_a, d_a, num_bytes, cudaMemcpyDeviceToHost );
```

Dummy Code Walkthrough 1



27

```
// walkthrough1.cu
#include <stdio.h>

int main()
{
    int dimx = 16;
    int num_bytes = dimx*sizeof(int);

    int *d_a=0, *h_a=0; // device and host pointers

    h_a = (int*)malloc(num_bytes);
    cudaMalloc( (void**)&d_a, num_bytes );

    if( 0==h_a || 0==d_a ) {
        printf("couldn't allocate memory\n");
        return 1;
    }

    cudaMemset( d_a, 0, num_bytes );

    //do something on the device to fill the buffer

    cudaMemcpy( h_a, d_a, num_bytes,
cudaMemcpyDeviceToHost );

    for(int i=0; i<dimx; i++)
        printf("%d ", h_a[i] );
    printf("\n");

    free( h_a );
    cudaFree( d_a );

    return 0;
}
```

Example: Shuffling Data



28

```
// Reorder values based on keys
// Each thread moves one element
__global__ void shuffle(int* prev_array, int*
    new_array, int* indices)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    new_array[i] = prev_array[indices[i]];
}
```

Host Code

```
int main()
{
    // Run grid of N/256 blocks of 256 threads
    each
```

```
    shuffle<<< N/256, 256>>>>(d_old, d_new, d_ind);
```

© 2008 NVIDIA Corporation

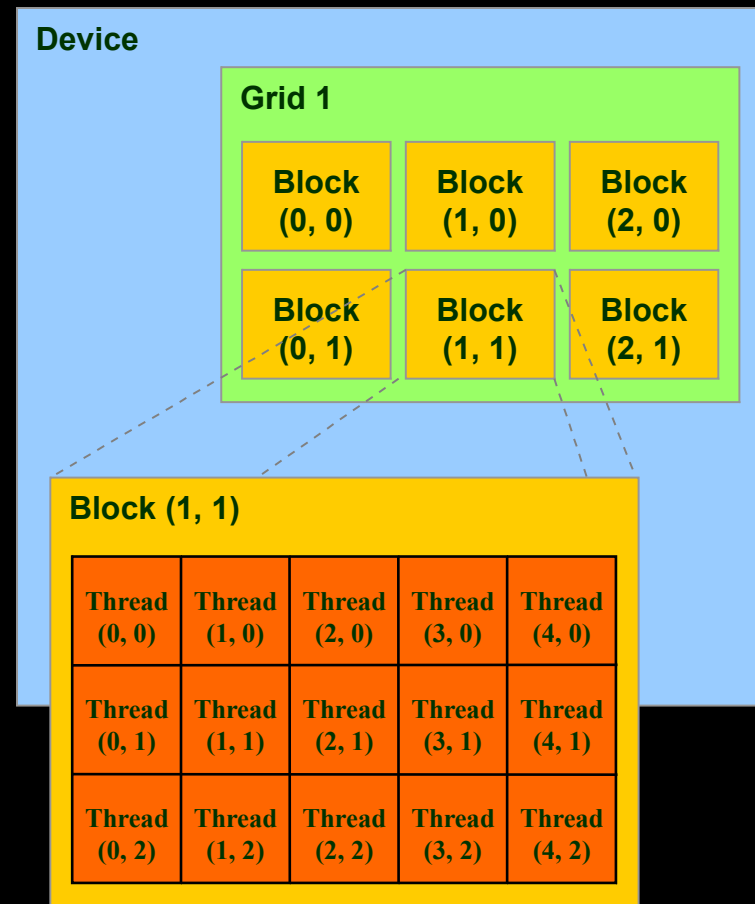
Τμήμα Ηλεκτρολόγων Μηχανικών & Μηχανικών
Υπολογιστών, Πανεπιστήμιο Θεσσαλίας

IDs and Dimensions



29

- **Threads:**
- 3D IDs, unique within a block
- **Blocks:**
- 2D IDs, unique within a grid
- **Dimensions set at launch**
- Can be unique for each grid
- **Built-in variables:**
- `threadIdx`, `blockIdx`
- `blockDim`, `gridDim`



Kernel with 2D Indexing



30

```
__global__ void kernel( int *a, int dimx, int dimy )
{
    int ix  = blockIdx.x*blockDim.x + threadIdx.x;
    int iy  = blockIdx.y*blockDim.y + threadIdx.y;
    int dimx = blockDim.x * gridDim.x;
    int idx = iy*dimx + ix;

    a[idx] = a[idx]+1;
}
```



31

```
int main() {  
    int dimx = 16;  
    int dimy = 16;  
    int num_bytes = dimx*dimy*sizeof(int);  
    int *d_a=0, *h_a=0; // device and host pointers
```

```
    h_a = (int*)malloc(num_bytes);  
    cudaMalloc( (void**)&d_a, num_bytes );
```

```
    if( 0==h_a || 0==d_a ) {  
        printf("couldn't allocate memory\n");  
        return 1;  
    }
```

```
    cudaMemset( d_a, 0, num_bytes );
```

```
    dim3 grid, block;  
    block.x = 4;  
    block.y = 4;  
    grid.x = dimx / block.x;  
    grid.y = dimy / block.y;
```

```
    kernel<<<grid, block>>>( d_a, dimx, dimy );
```

```
    cudaMemcpy( h_a, d_a, num_bytes, cudaMemcpyDeviceToHost );
```

```
    for(int row=0; row<dimy; row++) {  
        for(int col=0; col<dimx; col++)  
            printf("%d ", h_a[row*dimx+col] );  
        printf("\n");  
    }
```

```
    free( h_a );  
    cudaFree( d_a );
```

```
    return 0;
```

```
}
```

```
__global__ void kernel( int *a, int dimx, int dimy )  
{  
    int ix = blockIdx.x*blockDim.x + threadIdx.x;  
    int iy = blockIdx.y*blockDim.y + threadIdx.y;  
    int idx = iy*dimx + ix;  
  
    a[idx] = a[idx]+1;  
}
```

Blocks must be independent



32

- Any possible interleaving of blocks should be valid
- presumed to run to completion without pre-emption
- can run in any order
- can run concurrently OR sequentially
- Blocks may coordinate but not synchronize
- shared queue pointer: OK
- shared lock: BAD ... can easily deadlock
-
- Independence requirement gives scalability

A Simple Running Example

Matrix Multiplication



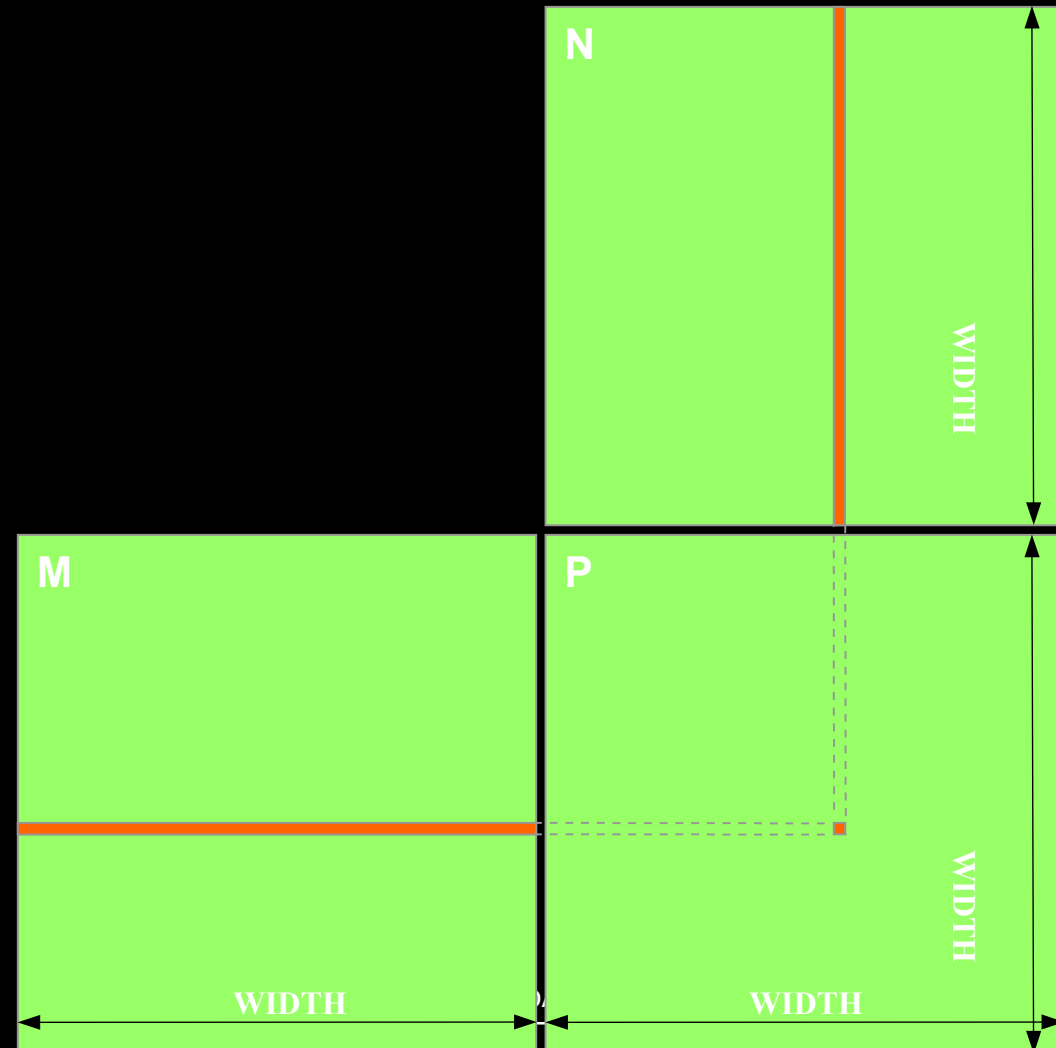
33

- A simple matrix multiplication example that illustrates the basic features of memory and thread management in CUDA programs
- Leave shared memory usage until later
- Local, register usage
- Thread ID usage
- Memory data transfer API between host and device
- Assume square matrix for simplicity



Programming Model: Square Matrix Multiplication Example

- $P = M * N$ of size **WIDTH x WIDTH**
- Without tiling:
- One **thread** calculates one element of P
- **M** and **N** are loaded **WIDTH times** from global memory



Memory Layout of a Matrix in C



35

$M_{0,0}$	$M_{1,0}$	$M_{2,0}$	$M_{3,0}$
$M_{0,1}$	$M_{1,1}$	$M_{2,1}$	$M_{3,1}$
$M_{0,2}$	$M_{1,2}$	$M_{2,2}$	$M_{3,2}$
$M_{0,3}$	$M_{1,3}$	$M_{2,3}$	$M_{3,3}$

M



$M_{0,0}$	$M_{1,0}$	$M_{2,0}$	$M_{3,0}$	$M_{0,1}$	$M_{1,1}$	$M_{2,1}$	$M_{3,1}$	$M_{0,2}$	$M_{1,2}$	$M_{2,2}$	$M_{3,2}$	$M_{0,3}$	$M_{1,3}$	$M_{2,3}$	$M_{3,3}$
-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------

Step 1: Matrix Multiplication

A Simple Host Version in C



36

// Matrix multiplication on the (CPU) host in double precision

```
void MatrixMulOnHost(float* M, float* N, float* P, int Width)
```

```
{
```

```
    for (int i = 0; i < Width; ++i)
```

```
        for (int j = 0; j < Width; ++j) {
```

```
            double sum = 0;
```

```
            for (int k = 0; k < Width; ++k) {
```

```
                double a = M[i * width + k];
```

```
                double b = N[k * width + j];
```

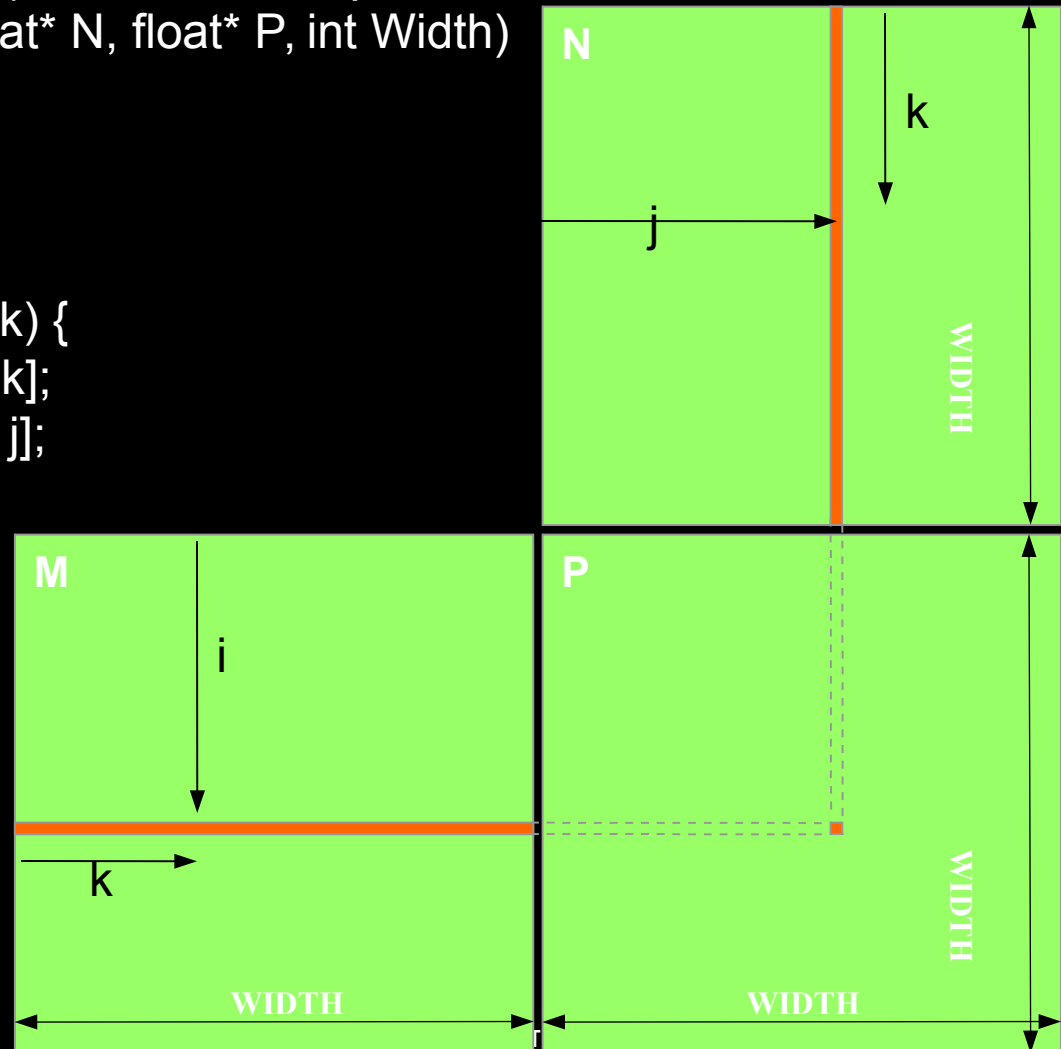
```
                sum += a * b;
```

```
            }
```

```
            P[i * Width + j] = sum;
```

```
        }
```

```
    }
```



Step 2: Input Matrix Data Transfer (Host-side Code)



37

```
void MatrixMulOnDevice(float* M, float* N, float* P, int Width)
{
    int size = Width * Width * sizeof(float);
    float* Md, Nd, Pd;
    ...
    // 1. Allocate and Load M, N to device memory
    cudaMalloc(&Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);

    cudaMalloc(&Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);

    // Allocate P on the device
    cudaMalloc(&Pd, size);
```

Step 3: Output Matrix Data Transfer (Host-side Code)



38

```
2. // Kernel invocation code – to be shown later
...

3. // Read P from the device
   cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);

   // Free device matrices
   cudaFree(Md); cudaFree(Nd); cudaFree (Pd);
}
```

Step 4: Kernel Function



39

// Matrix multiplication kernel – per thread code

```
__global__ void MatrixMulKernel(float* Md, float* Nd, float* Pd, int Width)
{
```

```
    // Pvalue is used to store the element of the matrix
```

```
    // that is computed by the thread
```

```
    float Pvalue = 0;
```

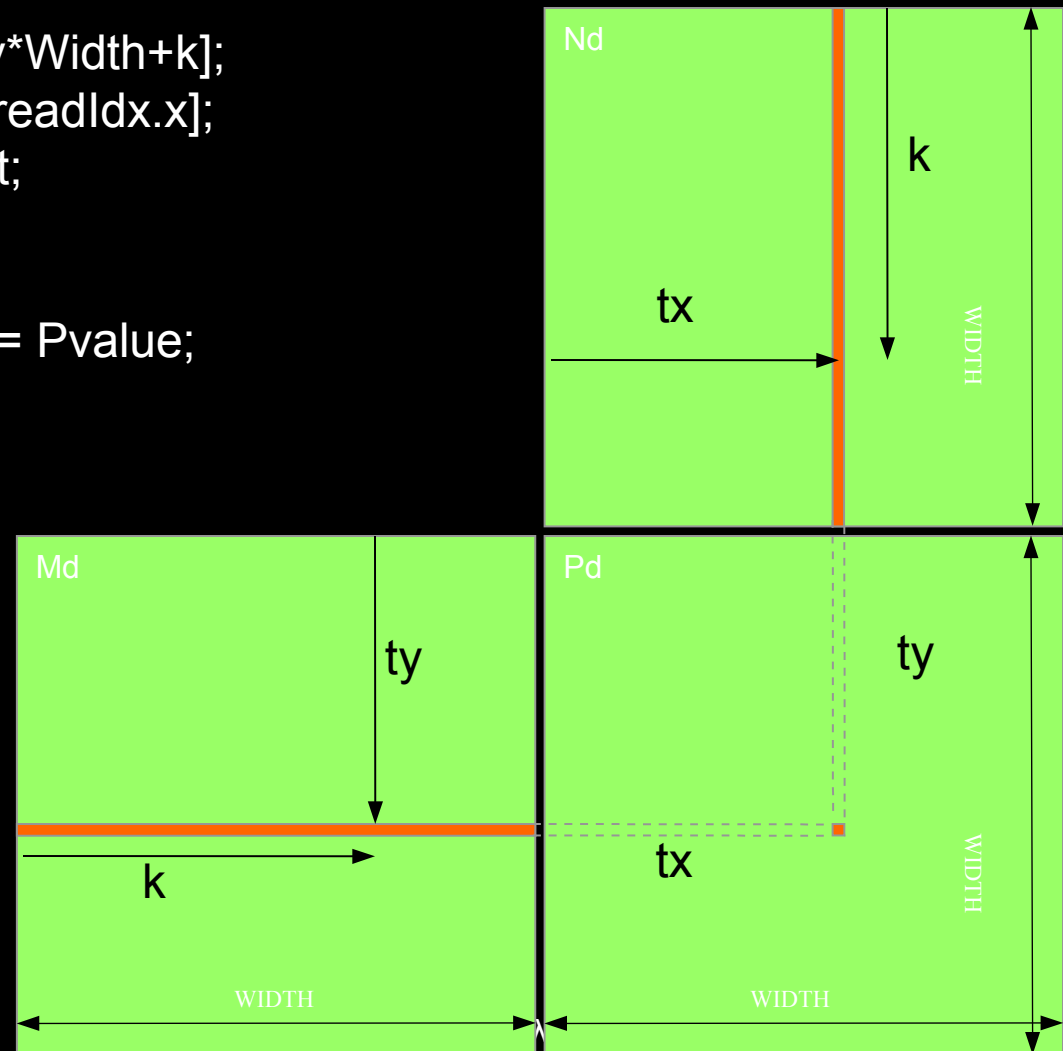
Step 4: Kernel Function (cont.)



40

```
for (int k = 0; k < Width; ++k) {  
    float Melement = Md[threadIdx.y*Width+k];  
    float Nelement = Nd[k*Width+threadIdx.x];  
    Pvalue += Melement * Nelement;  
}
```

```
Pd[threadIdx.y*Width+threadIdx.x] = Pvalue;  
}
```



Step 5: Kernel Invocation (Host-side Code)



41

```
// Setup the execution configuration  
dim3 dimGrid(1, 1);  
dim3 dimBlock(Width, Width);
```

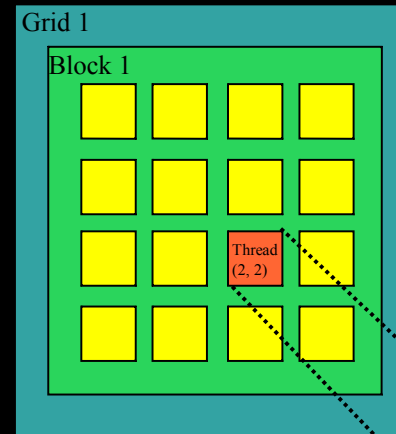
```
// Launch the device computation threads!  
MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);
```

Only One Thread Block Used

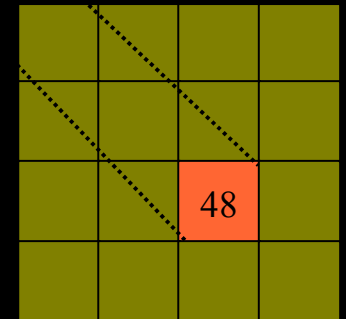
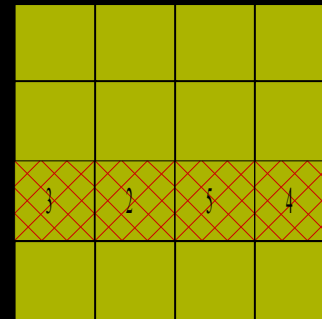
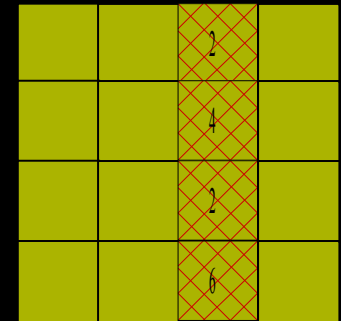


42

- One Block of threads compute matrix Pd
- Each thread computes one element of Pd
- Each thread
 - Loads a row of matrix Md
 - Loads a column of matrix Nd
 - Perform one multiply and addition for each pair of Md and Nd elements
 - Compute to off-chip memory access ratio close to 1:1 (not very high)
- Size of matrix limited by the number of threads allowed in a thread block



Nd



Md

Pd

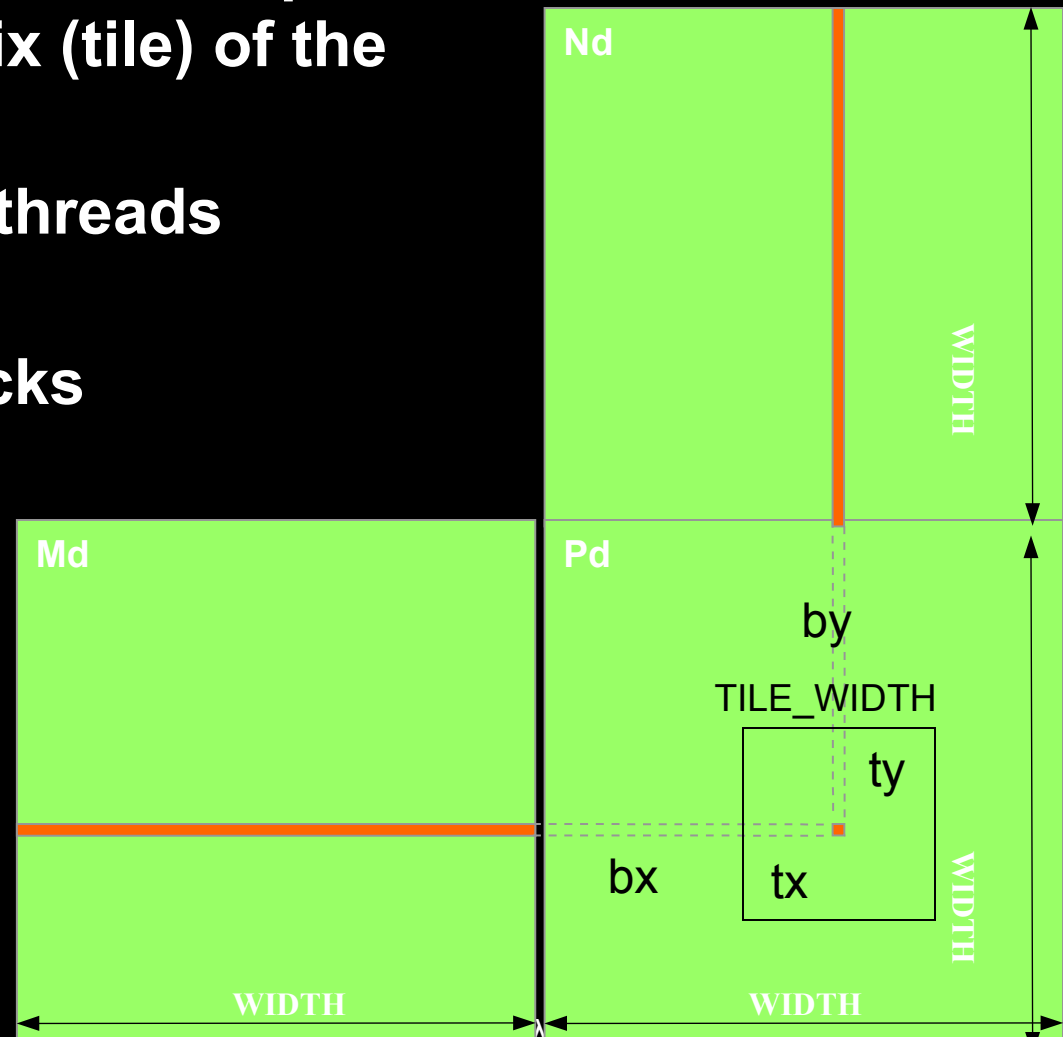
Step 7: Handling Arbitrary Sized Square Matrices (will cover later)



43

- Have each 2D thread block to compute a $(\text{TILE_WIDTH})^2$ sub-matrix (tile) of the result matrix
- Each has $(\text{TILE_WIDTH})^2$ threads
- Generate a 2D Grid of $(\text{WIDTH}/\text{TILE_WIDTH})^2$ blocks

You still need to put a loop around the kernel call for cases where $\text{WIDTH}/\text{TILE_WIDTH}$ is greater than max grid size (64K)!





USEFUL INFORMATION ON TOOLS

NVCC Compiler



45

- **NVIDIA provides a CUDA-C compiler**
 - **nvcc**
- **NVCC compiles device code then forwards code on to the host compiler (e.g. g++)**
- **Can be used to compile & link host only applications**

Example 1: Hello World



46

```
int main() {  
    printf("Hello World!\n");  
    return 0;  
}
```

Instructions:

1. Build and run the hello world code
2. Modify Makefile to use nvcc instead of g++
3. Rebuild and run

CUDA Example 1: Hello World



47

```
__global__ void mykernel(void) {  
  
    }  
  
int main(void) {  
    mykernel<<<1,1>>>();  
    printf("Hello World!\n");  
    return 0;  
}
```

Instructions:

1. Add kernel and kernel launch to main.cu
2. Try to build

CUDA Example 1: Build Considerations



48

- **Build failed**
 - Nvcc only parses .cu files for CUDA
- **Fixes:**
 - Rename main.cc to main.cu
 - OR
 - `nvcc -x cu`
 - Treat all input files as .cu files

Instructions:

1. Rename main.cc to main.cu
2. Rebuild and Run

Hello World! with Device Code



49

```
__global__ void mykernel(void) {  
}  
  
int main(void) {  
    mykernel<<<1,1>>>();  
    printf("Hello World!\n");  
    return 0;  
}
```

Output:

```
$ nvcc main.cu  
$ ./a.out  
Hello World!
```

mykernel (does nothing, somewhat anticlimactic!)

Developer Tools - Debuggers



50

NSIGHT



CUDA-GDB



CUDA MEMCHECK



NVIDIA Provided

allinea
DDT

TotalView®

3rd Party

<https://developer.nvidia.com/debugging-solutions>



Compiler Flags

- **Remember there are two compilers being used**
 - NVCC: Device code
 - Host Compiler: C/C++ code
- **NVCC supports some host compiler flags**
 - If flag is unsupported, use `-Xcompiler` to forward to host
 - e.g. `-Xcompiler -fopenmp`
- **Debugging Flags**
 - `-g`: Include host debugging symbols
 - `-G`: Include device debugging symbols
 - `-lineinfo`: Include line information

CUDA-MEMCHECK



52

- **Memory debugging tool**
 - No recompilation necessary
 - `%> cuda-memcheck ./exe`
- **Can detect the following errors**
 - Memory leaks
 - Memory errors (OOB, misaligned access, illegal instruction, etc)
 - Race conditions
 - Illegal Barriers
 - Uninitialized Memory
- **For line numbers use the following compiler flags:**
 - `-Xcompiler -rdynamic -lineinfo`

<http://docs.nvidia.com/cuda/cuda-memcheck>

Example 2: CUDA-MEMCHECK



53

Instructions:

1. Build & Run Example 2
Output should be the numbers 0-9
Do you get the correct results?
2. Run with cuda-memcheck
`%> cuda-memcheck ./a.out`
3. Add nvcc flags “-Xcompiler -rdynamic –lineinfo”
4. Rebuild & Run with cuda-memcheck
5. Fix the illegal write

<http://docs.nvidia.com/cuda/cuda-memcheck>

CUDA-GDB



54

- **cuda-gdb is an extension of GDB**
 - Provides seamless debugging of CUDA and CPU code
- **Works on Linux and Macintosh**
 - For a Windows debugger use NSIGHT Visual Studio Edition

<http://docs.nvidia.com/cuda/cuda-gdb>

Example 3: cuda-gdb



55

Instructions:

1. Run exercise 3 in cuda-gdb

```
%> cuda-gdb --args ./a.out
```

2. Run a few cuda-gdb commands:

```
(cuda-gdb) b main           //set break point at main
(cuda-gdb) r                 //run application
(cuda-gdb) l                 //print line context
(cuda-gdb) b foo             //break at kernel foo
(cuda-gdb) c                 //continue
(cuda-gdb) cuda thread       //print current thread
(cuda-gdb) cuda thread 10    //switch to thread 10
(cuda-gdb) cuda block        //print current block
(cuda-gdb) cuda block 1      //switch to block 1
(cuda-gdb) d                 //delete all break points
(cuda-gdb) set cuda memcheck on //turn on cuda
memcheck
(cuda-gdb) r                 //run from the beginning
```

3. Fix Bug

<http://docs.nvidia.com/cuda/cuda-gdb>

Developer Tools - Profilers

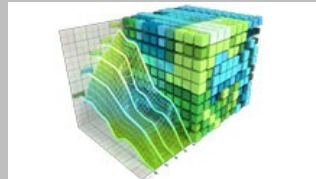


56

NSIGHT



NVVP

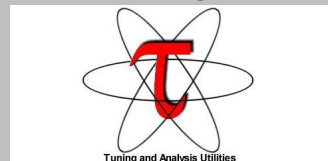


NVPROF

```
==20561== Profiling result:
Time(s)   Time    Calls      Avg      Min      Max   Name
49.88%  866.69ms  504798  1.7170us  1.5840us  2.0160us  void th
int, thrust::detail::device_generate_functor<thrust::detail::fill_
25.33%  440.05ms  252662  1.7410us  1.5360us  2.3680us  void th
c, thrust::detail::device_generate_functor<thrust::detail::fill_fo
17.07%  296.69ms   200    1.4830ms  1.2840ms  1.7253ms  kerComp
2.98%   51.81ms   200    259.89us  246.97us  264.83us  kerMake
1.18%   20.17ms   501    40.285us   928ns   17.677ms  [CUDA m
0.93%   16.19ms   200    80.991us   71.840us  90.751us  kerColl
0.73%   12.63ms   400    31.589us   14.720us  50.432us  [CUDA m
0.69%   12.07ms   200    60.376us   59.680us  62.304us  kerRema
0.63%   10.99ms   200    54.963us   52.680us  58.208us  kerMake
0.32%    5.552ms  200    27.761us   22.559us  33.152us  [CUDA w
0.12%    2.134ms   1     2.1342ms  2.1342ms  2.1342ms  void th
```

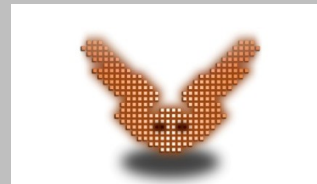
NVIDIA Provided

TAU



Tuning and Analysis Utilities

VampirTrace



3rd Party

<https://developer.nvidia.com/performance-analysis-tools>

NVPROF



57

Command Line Profiler

- **Compute time in each kernel**
- **Compute memory transfer time**
- **Collect metrics and events**
- **Support complex process hierarchy's**
- **Collect profiles for NVIDIA Visual Profiler**
- **No need to recompile**

Example 4: nvprof



58

Instructions:

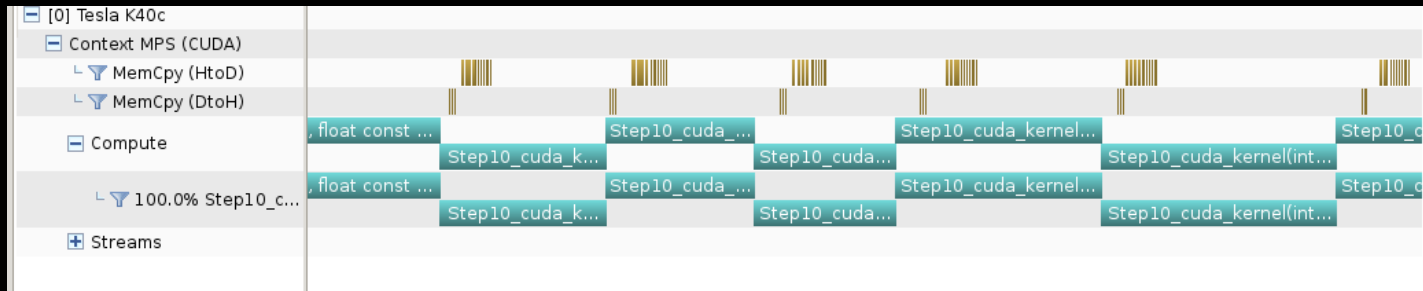
1. Collect profile information for the matrix add example
`%> nvprof ./a.out`
2. How much faster is add_v2 than add_v1?
3. View available metrics
`%> nvprof --query-metrics`
4. View global load/store efficiency
`%> nvprof --metrics
gld_efficiency,gst_efficiency ./a.out`
5. Store a timeline to load in NVVP
`%> nvprof -o profile.timeline ./a.out`
6. Store analysis metrics to load in NVVP
`%> nvprof -o profile.metrics --analysis-
metrics ./a.out`

NVIDIA's Visual Profiler (NVVP)



59

Timeline



Guided System

1. CUDA Application Analysis

2. Performance-Critical Kernels

3. Compute, Bandwidth, or Latency Bound

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results at right indicate that the performance of kernel "Step10_cuda_kernel" is most likely limited by compute.

[Perform Compute Analysis](#)

The most likely bottleneck to performance for this kernel is compute so you should first perform compute analysis to determine how it is limiting performance.

[Perform Latency Analysis](#)

[Perform Memory Bandwidth Analysis](#)

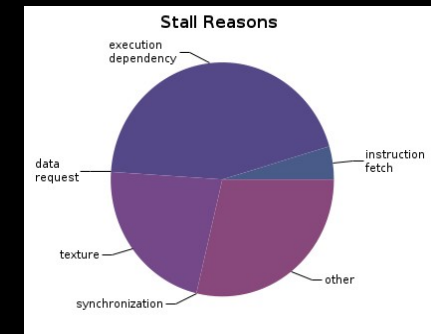
Instruction and memory latency and memory bandwidth are likely not the primary performance bottlenecks for this kernel, but you may still want to perform those analyses.

[Rerun Analysis](#)

If you modify the kernel you need to rerun your application to update this analysis.

Analysis

CUDA Shared Memory		
Local Loads	0	0 B/s
Local Stores	0	0 B/s
Global Loads	0	0 B/s
Global Stores	0	0 B/s
Global Loads	0	0 B/s
Global Stores	0	0 B/s
L1/Shared total	0	0 B/s
L2 Cache		
Reads	6339426	236.738 GB/s
Writes	31614	1.173 GB/s
Total	6370840	237.912 GB/s
Texture Cache		
Reads	6450496	240.886 GB/s
Device Memory		
Reads	1562634	58.355 GB/s
Writes	7504	280.228 MB/s
Total	1570138	58.635 GB/s
System Memory [PCIe configuration: Gen3 x16, 8 Gbit/s]		
Reads	0	0 B/s
Writes	4	149.375 kB/s
Total	4	149.375 kB/s



Example 4: NVVP



60

Instructions:

1. Import nvprof profile into NVVP

Launch nvvp

Click File/ Import/ Nvprof/ Next/ Single process/
Next / Browse

Select profile.timeline

Add Metrics to timeline

Click on 2nd Browse

Select profile.metrics

Click Finish

2. Explore Timeline

Control + mouse drag in timeline to zoom in

Control + mouse drag in measure bar (on top) to
measure time

Example 4: NVVP

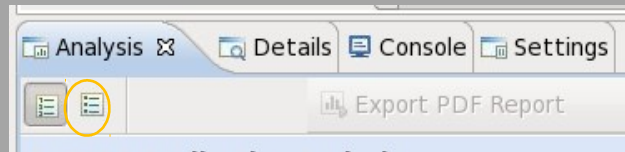


61

Instructions:

1. Click on a kernel
2. On Analysis tab click on the unguided analysis
- 3.

1.



2. Click Analyze All
Explore metrics and properties
What differences do you see between the two kernels?

Note:

If kernel order is non-deterministic you can only load the timeline or the metrics but not both.

If you load just metrics the timeline looks odd but metrics are correct.

Example 4: NVVP



62

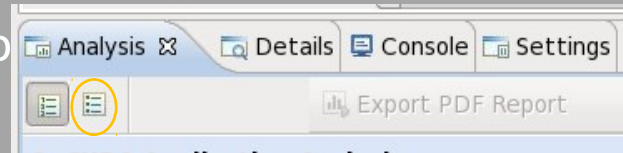
Let's now generate the same data within NVVP

Instructions:

1. Click File / New Session / Browse

Select Example 4/a.o

Click Next / Finish



2. Click on a kernel

Select Unguided Analysis

Click Analyze All



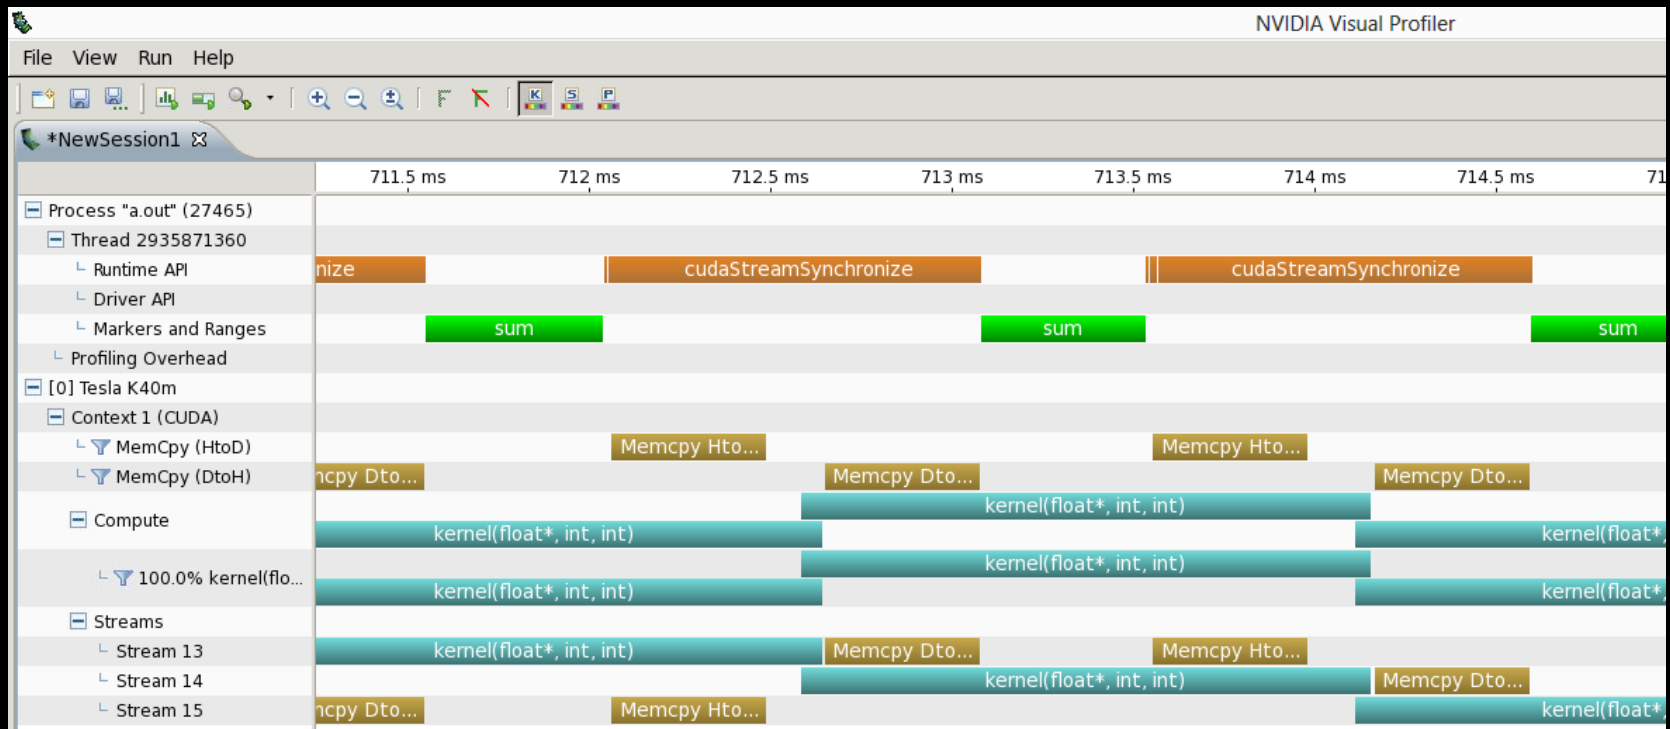
- **Our current tools only profile API calls on the host**
 - What if we want to understand better what the host is doing?
- **The NVTX library allows us to annotate profiles with ranges**
 - Add: `#include <nvToolsExt.h>`
 - Link with: `-lnvToolsExt`
- **Mark the start of a range**
 - `nvtxRangePushA("description");`
- **Mark the end of a range**
 - `nvtxRangePop();`
- **Ranges are allowed to overlap**

<http://devblogs.nvidia.com/parallelforall/cuda-pro-tip-generate-custom-application-profile-timelines-nvtx/>

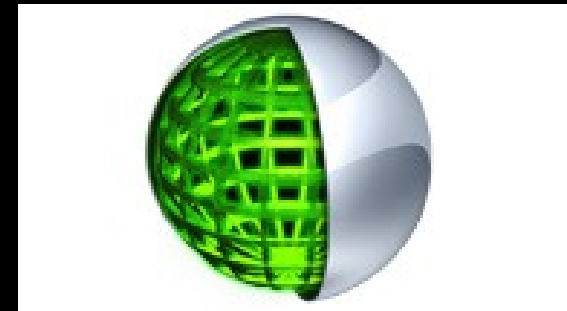
NVTX Profile



64



- **CUDA enabled Integrated Development Environment**
 - Source code editor: syntax highlighting, code refactoring, etc
 - Build Manger
 - Visual Debugger
 - Visual Profiler
- **Linux/Macintosh**
 - Editor = Eclipse
 - Debugger = cuda-gdb with a visual wrapper
 - Profiler = NVVP
- **Windows**
 - Integrates directly into Visual Studio
 - Profiler is NSIGHT VSE



Example 4: NSIGHT



66

Let's import an existing Makefile project into NSIGHT

Instructions:

1. Run nsight
Select default workspace
2. Click File / New / Makefile Project With Existing CodeTest
3. Enter Project Name and select the Example15 directory
4. Click Finish
5. Right Click On Project / Properties / Run Settings / New / C++ Application
6. Browse for Example 4/a.out
7. In Project Explorer double click on main.cu and explore source
8. Click on the build icon
9. Click on the run icon
10. Click on the profile icon



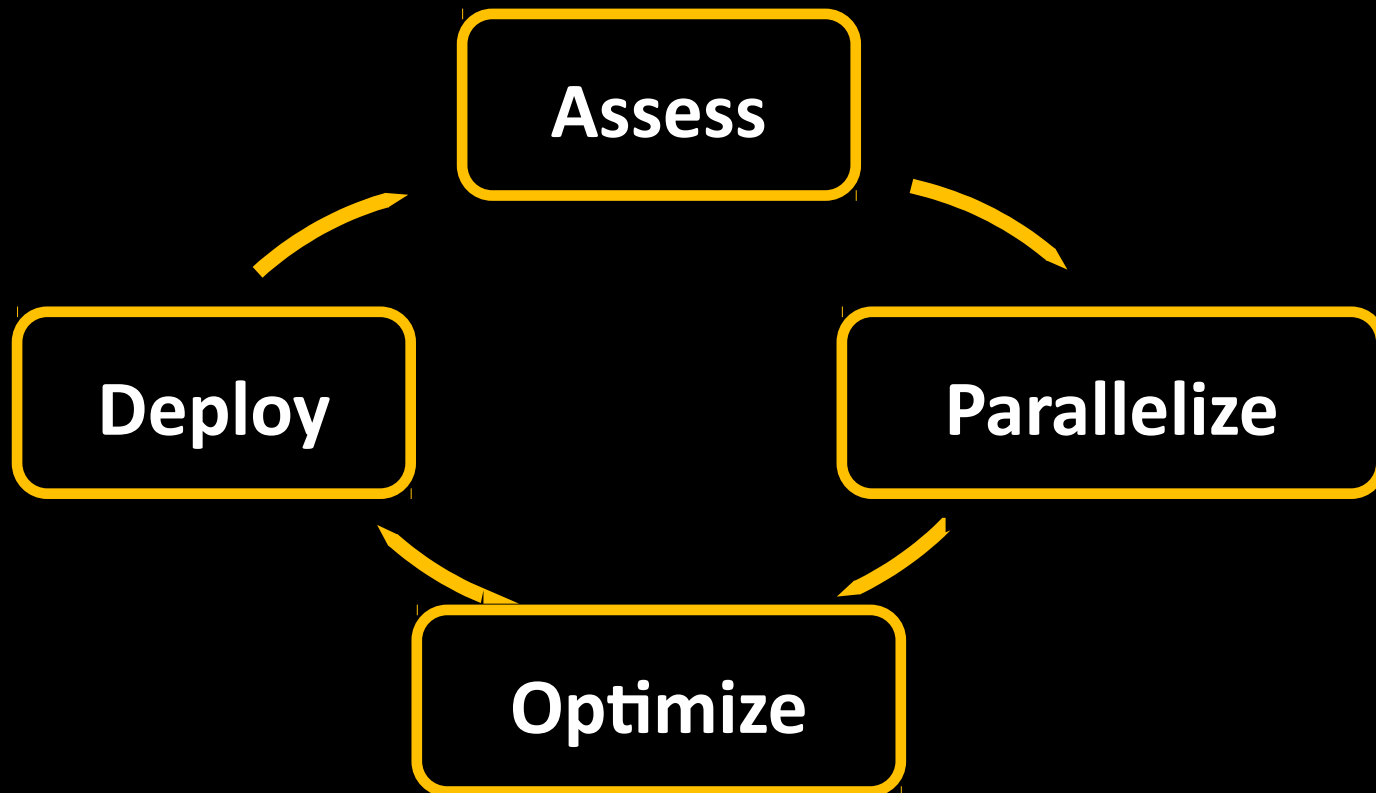
Profiler Summary

- **Many profile tools are available**
- **NVIDIA Provided**
 - NVPROF: Command Line
 - NVVP: Visual profiler
 - NSIGHT: IDE (Visual Studio and Eclipse)
- **3rd Party**
 - TAU
 - VAMPIR

Optimization



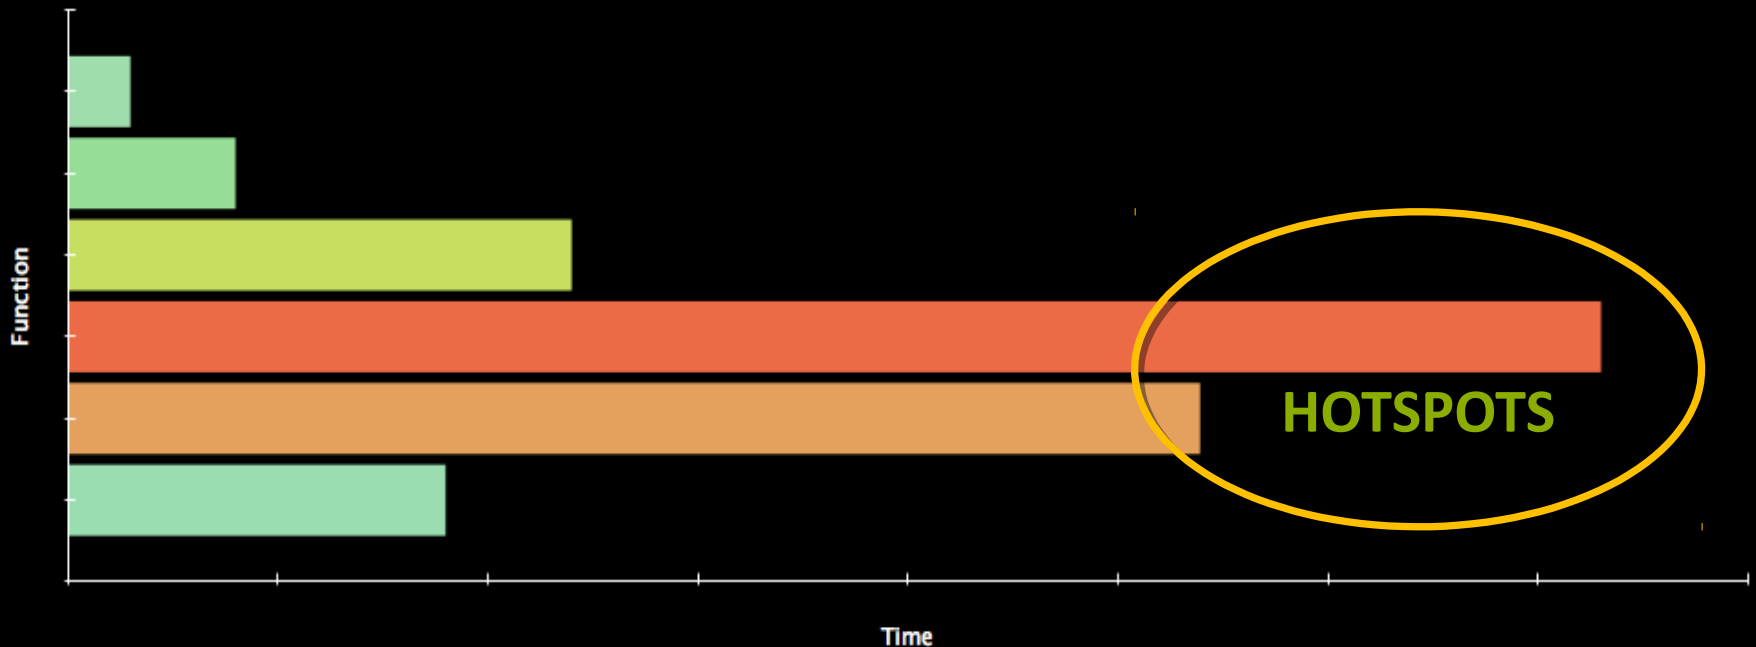
68



Assess



69



- Profile the code, find the hotspot(s)
- Focus your attention where it will give the most benefit

Parallelize



70

Applications

Libraries

Compiler
Directives

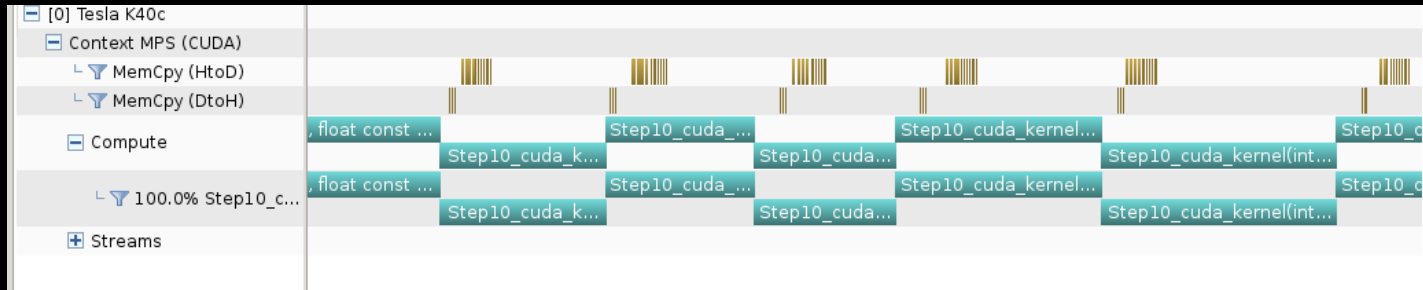
Programming
Languages

Optimize



71

Timeline



Guided System

1. CUDA Application Analysis

2. Performance-Critical Kernels

3. Compute, Bandwidth, or Latency Bound

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results at right indicate that the performance of kernel "Step10_cuda_kernel" is most likely limited by compute.

[Perform Compute Analysis](#)

The most likely bottleneck to performance for this kernel is compute so you should first perform compute analysis to determine how it is limiting performance.

[Perform Latency Analysis](#)

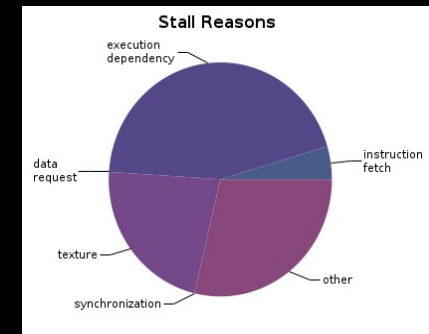
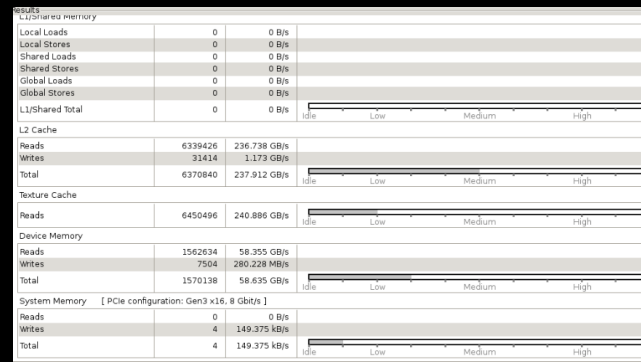
[Perform Memory Bandwidth Analysis](#)

Instruction and memory latency and memory bandwidth are likely not the primary performance bottlenecks for this kernel, but you may still want to perform those analyses.

[Rerun Analysis](#)

If you modify the kernel you need to rerun your application to update this analysis.

Analysis



Bottleneck Analysis



72

- Don't assume an optimization was wrong
- Verify if it was wrong with the profiler

129 GB/s ➡ 84 GB/s

L1/Shared Memory			
Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Shared Loads	2097152	1,351.979 GB/s	
Shared Stores	131072	84.499 GB/s	
Global Loads	131072	42.249 GB/s	
Global Stores	131072	42.249 GB/s	
Atomic	0	0 B/s	
L1/Shared Total	2490368	1,520.977 GB/s	

gpuTranspose_kernel(int, int, float const *, float*)	
Start	547.303 ms (5)
End	547.716 ms (5)
Duration	413.872 μs
Grid Size	[64,64,1]
Block Size	[32,32,1]
Registers/Thread	10
Shared Memory/Block	4 KiB
Efficiency	
Global Load Efficiency	100%
Global Store Efficiency	100%
Shared Efficiency	5.9%
Warp Execution Efficiency	100%
Non-Predicated Warp Execution Efficiency	97.1%
Occupancy	
Achieved	86.7%
Theoretical	100%
Shared Memory Configuration	
Shared Memory Requested	48 KiB
Shared Memory Executed	48 KiB

Shared Memory Alignment and Access Pattern

Memory bandwidth is used most efficiently when each shared memory load and store has proper alignment and access pattern.

Optimization: Select each entry below to open the source code to a shared load or store within the kernel with an inefficient alignment or access pattern. For each access pattern of the memory access.

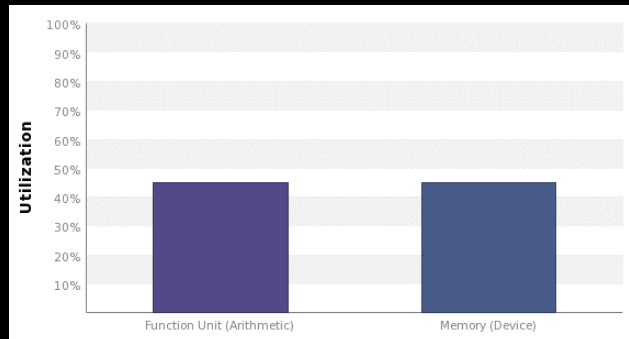
Line / File	main.cu - /home/jluitjens/code/CudaHandsOn/Example19
49	Shared Load Transactions/Access = 16, Ideal Transactions/Access = 1 [2097152 transactions for 131072 total executions]

Performance Analysis



73

gpuTranspose_kernel(int, int, float const *, float*)	
Start	770.067
End	770.324
Duration	256.714
Grid Size	[64,64,1
Block Size	[32,32,1
Registers/Thread	10
Shared Memory/Block	4.125 KiB
Efficiency	
Global Load Efficiency	100%
Global Store Efficiency	100%
Shared Efficiency	50%
Warp Execution Efficiency	100%
Non-Predicated Warp Execution Efficiency	97.1%
Occupancy	
Achieved	87.7%
Theoretical	100%
Shared Memory Configuration	
Shared Memory Requested	48 KiB
Shared Memory Executed	48 KiB



84 GB/s → 137 GB/s

L1/Shared Memory			
Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Shared Loads	131072	138.433 GB/s	
Shared Stores	131720	139.118 GB/s	
Global Loads	131072	69.217 GB/s	
Global Stores	131072	69.217 GB/s	
Atomic	0	0 B/s	
L1/Shared Total	524936	415.984 GB/s	
L2 Cache			
L1 Reads	524288	69.217 GB/s	
L1 Writes	524288	69.217 GB/s	
Texture Reads	0	0 B/s	
Atomic	0	0 B/s	
Noncoherent Reads	0	0 B/s	
Total	1048576	138.433 GB/s	
Texture Cache			
Reads	0	0 B/s	
Device Memory			
Reads	524968	69.306 GB/s	
Writes	524289	69.217 GB/s	
Total	1049257	138.523 GB/s	



- **Results of floating-point computations will slightly differ on some GPUs because of:**
 - Different compiler outputs, instruction sets
 - Use of extended precision for intermediate results
- There are various options to force strict single precision on the host



75

ATOMICS

The Problem



76

- **How do you do global communication?**
- **Finish a grid and start a new one**
-
-



- Finish a kernel and start a new one
- All writes from all threads complete before a kernel finishes
- If kernel invocations are synchronous.
- This is not the case with Fermi and later!

```
step1<<<grid1,blk1>>>(...);  
// The system ensures that all  
// writes from step1  
complete.  
step2<<<grid2,blk2>>>(...);
```

Global Communication



78

- **Would need to decompose kernels into before and after parts**

Race Conditions



79

- Or, write to a predefined memory location
- Race condition! Updates can be lost

Race Conditions



80

```
threadId:0
threadId:1917
    // vector[0] was equal to 0
vector[0] += 5;                vector[0] += 1;
...                            ...
a = vector[0];                a = vector[0];
```

- What is the value of a in thread 0?
- What is the value of a in thread 1917?

Race Conditions



81

- Thread 0 could have finished execution before 1917 started
- Or the other way around
- Or both are executing at the same time

Race Conditions



82

- **Answer: not defined by the programming model, can be arbitrary**

Atomics



83

- CUDA provides **atomic** operations to deal with this problem



Atomics



84

- An atomic operation guarantees that only a single thread has access to a piece of memory while an operation completes
- The name atomic comes from the fact that it is uninterruptable
- No dropped data, but ordering is still arbitrary
- Different types of atomic instructions
- `atomic{Add, Sub, Exch, Min, Max, Inc, Dec, CAS, And, Or, Xor}`
- More types in Fermi and later
-
-

Example: Histogram



85

```
// Determine frequency of colors in a picture
// colors have already been converted into ints
// Each thread looks at one pixel and increments
// a counter atomically
__global__ void histogram(int* color,
                          int* buckets)
{
    int i = threadIdx.x
          + blockDim.x * blockIdx.x;
    int c = colors[i];
    atomicAdd(&buckets[c], 1);
}
```

Example: Workqueue



86

```
// For algorithms where the amount of work per  
    item  
// is highly non-uniform, it often makes sense for  
// to continuously grab work from a queue
```

```
__global__
```

```
void workq(int* work_q, int* q_counter,  
           int* output, int queue_max)
```

```
{
```

```
    int i = threadIdx.x  
           + blockDim.x * blockIdx.x;
```

```
    int q_index =
```

```
        atomicInc(q_counter, queue_max);
```

```
    int result = do_work(work_q[q_index]);
```

```
    output[i] = result;
```

Atomics



87

- **Atomics are slower than normal load/store**
- **You can have the whole machine queuing on a single location in memory**
- **Atomics unavailable on older (really old now) architectures!**

Atomics



88

- **atomicAdd** returns the previous value at a certain address
- Useful for grabbing variable amounts of data from a list

Example: Global Min/Max (Naive)



```
// If you require the maximum across all threads 89
// in a grid, you could do it with a single global
// maximum value, but it will be VERY slow
```

__global__

```
void global_max(int* values, int* gl_max)
{
    int i = threadIdx.x
        + blockDim.x * blockIdx.x;
    int val = values[i];
    atomicMax(gl_max, val);
}
```

Example: Global Min/Max (Better)



```
// introduce intermediate maximum results, so that 90
// most threads do not try to update the global
max
```

```
__global__
```

```
void global_max(int* values, int* max,
                int *reg_max,
                int num_regions)
```

```
{
```

```
    // i and val as before ...
```

```
    int region = i % num_regions;
```

```
    if(atomicMax(&reg_max[region], val) <
val)
```

```
    {
```

```
        atomicMax(max, val);
```

© 2008 NVIDIA Corporation

Global Min/Max



91

- **Single value causes serial bottleneck**
- **Create hierarchy of values for more parallelism**
- **Performance will still be slow, so use judiciously**
- **See next lecture for even better version!**

Summary



92

- Can't use normal load/store for inter-thread communication because of **race conditions**
-
- Use **atomic instructions** for sparse and/or unpredictable global communication
- See next lectures for shared memory and scan for other communication patterns
-
- **Decompose data** (very limited use of single global sum/max/min/etc.) for more parallelism
-



SM EXECUTION & DIVERGENCE

How an SM executes threads



94

- Overview of how a Stream Multiprocessor works
- SIMT Execution
- Divergence

Scheduling Blocks onto SMs



95

Streaming Multiprocessor



Thread Block 5

Thread Block 27

Thread Block 61

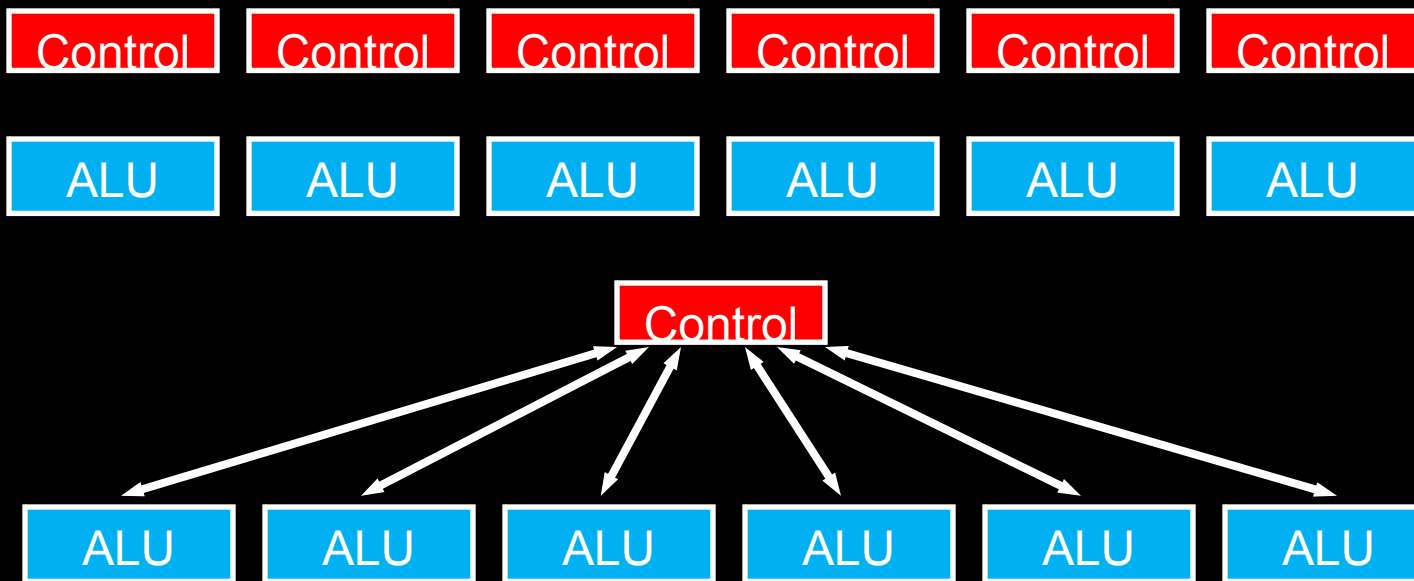
Thread Block 2001

- **HW Schedules thread blocks onto available SMs**
- **No guarantee of ordering among thread blocks**
- **HW will schedule thread blocks as soon as a previous thread block finishes**

Warps



96



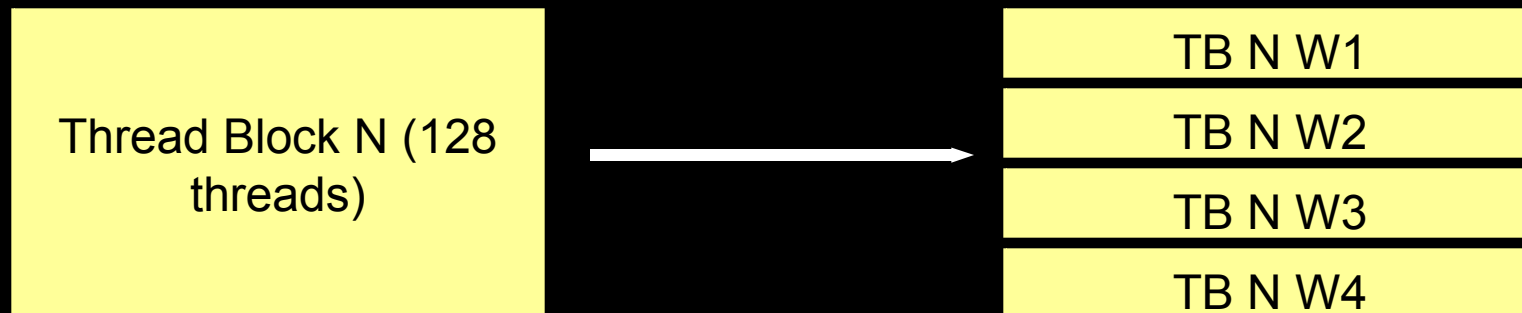
- A **warp** = 32 threads launched together
- Usually, execute together as well

Mapping of Thread Blocks



97

- Each thread block is mapped to one or more warps
- The hardware schedules each warp independently

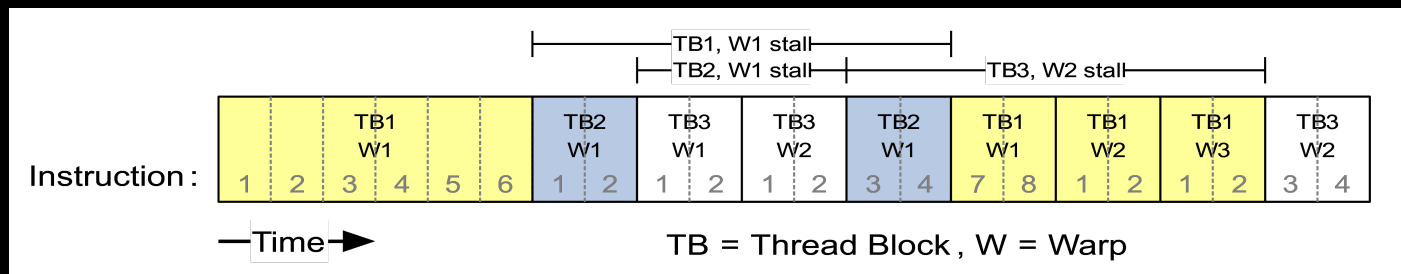


Thread Scheduling Example



98

- SM implements zero-overhead warp scheduling
- At any time, only one of the warps is executed by SM *
- Warps whose next instruction has its inputs ready for consumption are eligible for execution
- Eligible Warps are selected for execution on a prioritized scheduling policy
- All threads in a warp execute the same instruction when selected



Control Flow Divergence



99

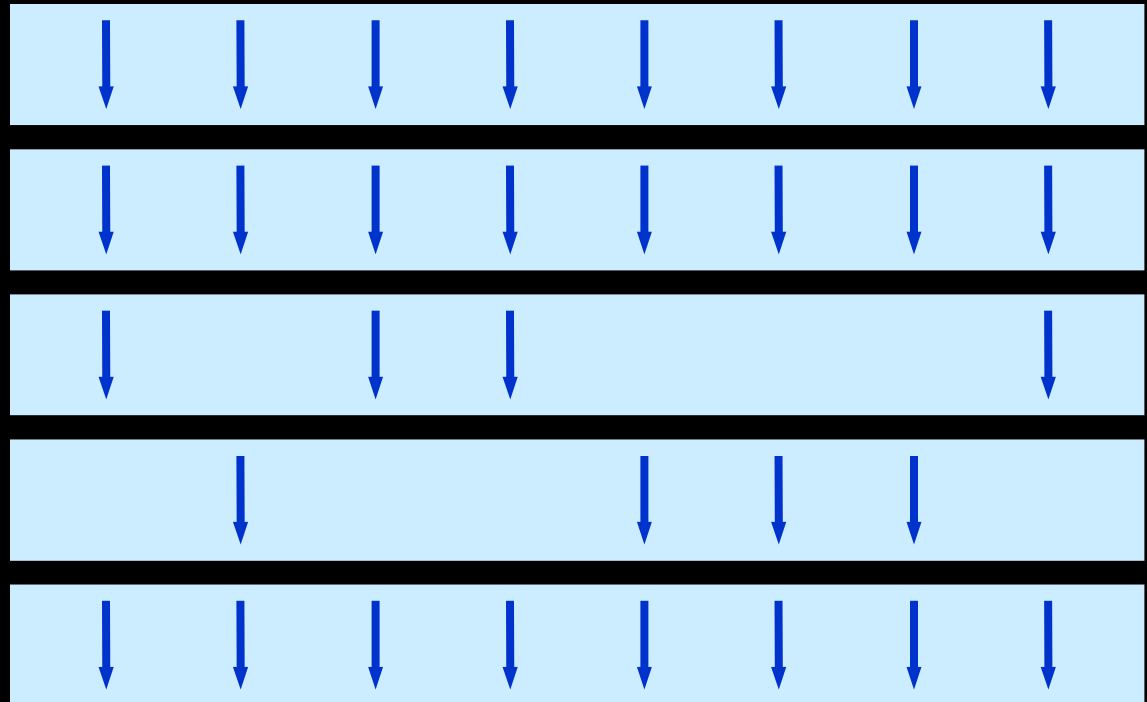
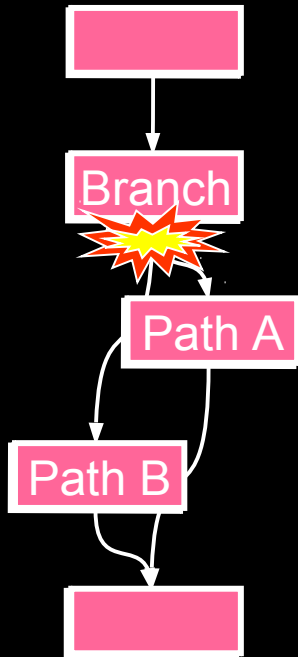
What happens if you have the following code?

```
if (foo (threadIdx.x) )  
{  
    do_A () ;  
}  
else  
{  
    do_B () ;  
}
```

Control Flow Divergence



100



From Fung et al. MICRO '07

Control Flow Divergence



101

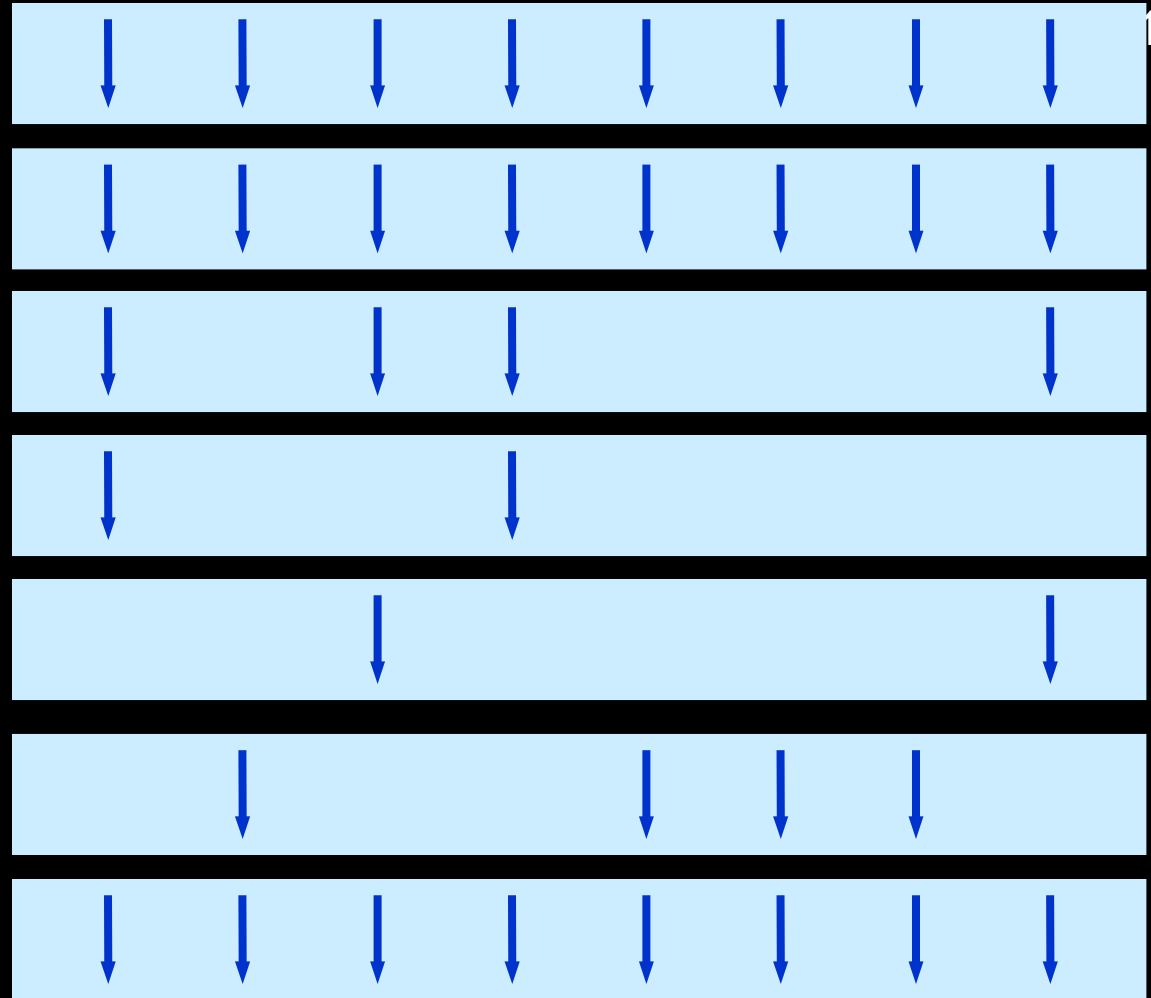
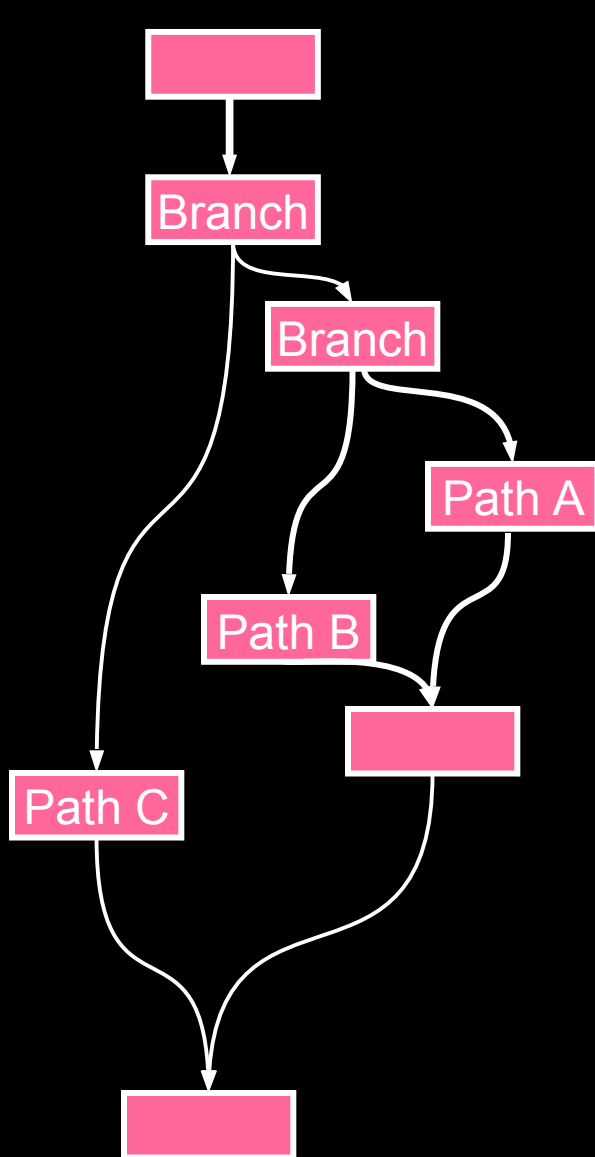
- Nested branches are handled as well

```
if (foo (threadIdx.x) )  
{  
    if (bar (threadIdx.x) )  
        do_A () ;  
    else  
        do_B () ;  
}  
else  
    do_C () ;
```

Control Flow Divergence



102



Control Flow Divergence



103

- You don't have to worry about divergence for correctness (*)
- You might have to think about it for performance
- Depends on your branch conditions

Control Flow Divergence



104

- Performance drops off with the degree of divergence

-
-

```
● switch (threadIdx.x % N)
● {
●     case 0:
●         ...
●     case 1:
●         ...
● }
```


Divergence



105

